Dynamic Memory Allocation: Advanced Concepts
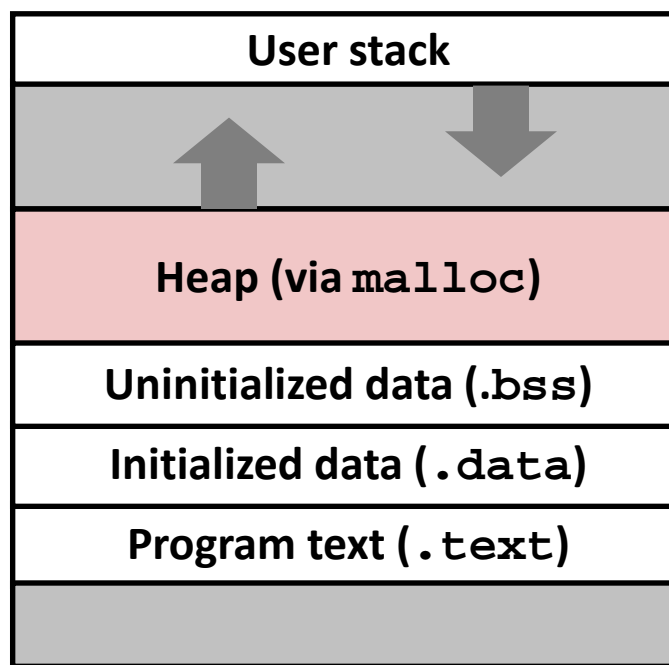
15-213: Introduction to Computer Systems
20th Lecture, April 4, 2017

Instructors:

Seth Copen Goldstein & Franz Franchetti

Dynamic Memory Allocation

- **Programmers use *dynamic memory allocators* (such as `malloc`) to acquire VM at run time.**
  - For data structures whose size is only known at runtime.

- **Dynamic memory allocators manage an area of process virtual memory known as the *heap*.**

| Application |
|---|
| **Dynamic Memory Allocator** |
| **Heap** |

User stack
Heap (via `malloc`)
Uninitialized data (`.bss`)
Initialized data (`.data`)
Program text (`.text`)

Top of heap (`brk ptr`)

0

Last Lecture: Keeping Track of Free Blocks

■ **Method 1: *Implicit list* using length—links all blocks**

■ **Method 2: *Explicit list* among the free blocks using pointers**

■ **Method 3: *Segregated free list***

- Different free lists for different size classes

■ **Method 4: *Blocks sorted by size***

- Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

Summary: Implicit Lists

- **Implementation: very simple**

- **Allocate cost:**
  - linear time worst case

- **Free cost:**
  - constant time worst case
  - even with coalescing

- **Memory usage:**
  - will depend on placement policy
  - First-fit, next-fit or best-fit

- **Not used in practice for `malloc/free` because of linear-time allocation**
  - used in many special purpose applications

- **However, the concepts of splitting and boundary tag coalescing are general to *all* allocators**

Today

- **Explicit free lists**
- **Segregated free lists**
- **Garbage collection**
- **Memory-related perils and pitfalls**

Keeping Track of Free Blocks

- **Method 1: *Implicit free list* using length—links all blocks**

| 5 | | | | | 4 | | | | 6 | | | | | 2 | |

- **Method 2: *Explicit free list* among the free blocks using pointers**

| 5 | | | | 4 | | | 6 | | | | | 2 | |

- **Method 3: *Segregated free list***
  - Different free lists for different size classes

- **Method 4: *Blocks sorted by size***
  - Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

Explicit Free Lists

Allocated (as before)

Size	a
Payload and padding	
Size	a

Free

Size	a
Next	
Prev	
Size	a

- **Maintain list(s) of *free* blocks, not *all* blocks**
  - The "next" free block could be anywhere
    - So we need to store forward/back pointers, not just sizes
  - Still need boundary tags for coalescing
  - Luckily we track only free blocks, so we can use payload area

Explicit Free Lists

- **Logically:**

- **Physically: blocks can be in any order**

Forward (next) links

Back (prev) links

Allocating From Explicit Free Lists

conceptual graphic

Before

After   *(with splitting)*

● = malloc(…)

Freeing With Explicit Free Lists

- *Insertion policy*: **Where in the free list do you put a newly freed block?**

- **Unordered**
  - LIFO (last-in-fir...
    - Insert free...
  - FIFO (first-in-fi...
    - Insert free...
  - *Pro:* simple an...
  - *Con:* studies su...

- **Address-ordere...**
  - Insert freed blo...
    $$addr(prev) < addr(curr) < addr(next)$$
  - *Con:* requires search
  - *Pro:* studies suggest fragmentation is lower than LIFO/FIFO

Aside: Premature Optimization

Don't!

Freeing With Explicit Free Lists

- *Insertion policy*: **Where in the free list do you put a newly freed block?**
- **Unordered**
  - LIFO (last-in-first-out) policy
    - Insert freed block at the beginning of the free list
  - FIFO (first-in-first-out) policy
    - Insert freed block at the end of the free list
  - *Pro:* simple and constant time
  - *Con:* studies suggest fragmentation is worse than address ordered
- **Address-ordered policy**
  - Insert freed blocks so that free list blocks are always in address order:
    addr(prev) < addr(curr) < addr(next)
  - *Con:* requires search
  - *Pro:* studies suggest fragmentation is lower than LIFO/FIFO

Freeing With a LIFO Policy (Case 1)

conceptual graphic

Before

free(●)

Root

- **Insert the freed block at the root of the list**

After

Root

Freeing With a LIFO Policy (Case 2)

conceptual graphic

Before

free(●)

Root

- **Splice out successor block, coalesce both memory blocks and insert the new block at the root of the list**

After

Root

Freeing With a LIFO Policy (Case 3)

conceptual graphic

Before

free(●)

Root

- **Splice out predecessor block, coalesce both memory blocks, and insert the new block at the root of the list**

After

Root

Freeing With a LIFO Policy (Case 4)

conceptual graphic

- **Splice out predecessor and successor blocks, coalesce all 3 memory blocks and insert the new block at the root of the list**

Some Advice: An Implementation Trick

- **Use circular, doubly-linked list**

- **Support multiple approaches with single data structure**

- **First-fit vs. next-fit**
  - Either keep free pointer fixed or move as search list

- **LIFO vs. FIFO**
  - Insert as next block (LIFO), or previous block (FIFO)

Explicit List Summary

- **Comparison to implicit list:**
  - Allocate is linear time in number of *free* blocks instead of *all* blocks
    - *Much faster* when most of the memory is full
  - Slightly more complicated allocate and free since needs to splice blocks in and out of the list
  - Some extra space for the links (2 extra words needed for each block)
    - Does this increase internal fragmentation?

- **Most common use of linked lists is in conjunction with segregated free lists**
  - Keep multiple linked lists of different size classes, or possibly for different types of objects

Today

- **Explicit free lists**
- **Segregated free lists**
- **Garbage collection**
- **Memory-related perils and pitfalls**

Segregated List (Seglist) Allocators

- **Each *size class* of blocks has its own free list**

- **Often have separate classes for each small size**
- **For larger sizes: One class for each two-power size**

Seglist Allocator

■ **Given an array of free lists, each one for some size class**

■ **To allocate a block of size *n*:**
   ▪ Search appropriate free list for block of size *m > n*
   ▪ If an appropriate block is found:
      ▪ Split block and place fragment on appropriate list (optional)
   ▪ If no block is found, try next larger class
   ▪ Repeat until block is found

■ **If no block is found:**
   ▪ Request additional heap memory from OS (using `sbrk()`)
   ▪ Allocate block of *n* bytes from this new memory
   ▪ Place remainder as a single free block in largest size class.

Seglist Allocator (cont.)

- **To free a block:**
  - Coalesce and place on appropriate list

- **Advantages of seglist allocators**
  - Higher throughput
    - log time for power-of-two size classes
  - Better memory utilization
    - First-fit search of segregated free list approximates a best-fit search of entire heap.
    - Extreme case: Giving each block its own size class is equivalent to best-fit.

More Info on Allocators

- **D. Knuth, "*The Art of Computer Programming*", 2nd edition, Addison Wesley, 1973**
    - The classic reference on dynamic storage allocation

- **Wilson et al, "*Dynamic Storage Allocation: A Survey and Critical Review*", Proc. 1995 Int'l Workshop on Memory Management, Kinross, Scotland, Sept, 1995.**
    - Comprehensive survey
    - Available from CS:APP student site (csapp.cs.cmu.edu)

Today

- **Explicit free lists**

- **Segregated free lists**

- **Garbage collection**

- **Memory-related perils and pitfalls**

Implicit Memory Management: Garbage Collection

■ *Garbage collection:* **automatic reclamation of heap-allocated storage—application never has to free**

```
void foo() {
    int *p = malloc(128);
    return; /* p block is now garbage */
}
```

■ **Common in many dynamic languages:**

- Python, Ruby, Java, Perl, ML, Lisp, Mathematica

■ **Variants ("conservative" garbage collectors) exist for C and C++**

- However, cannot necessarily collect all garbage

Garbage Collection

- **How does the memory manager know when memory can be freed?**
  - In general we cannot know what is going to be used in the future since it depends on conditionals
  - But we can tell that certain blocks cannot be used if there are no pointers to them

- **Must make certain assumptions about pointers**
  - Memory manager can distinguish pointers from non-pointers
  - All pointers point to the start of a block
  - Cannot hide pointers
    (e.g., by coercing them to an `int`, and then back again)

Classical GC Algorithms

- **Mark-and-sweep collection (McCarthy, 1960)**
  - Does not move blocks (unless you also "compact")
- **Reference counting (Collins, 1960)**
  - Does not move blocks (not discussed)
- **Copying collection (Minsky, 1963)**
  - Moves blocks (not discussed)
- **Generational Collectors (Lieberman and Hewitt, 1983)**
  - Collection based on lifetimes
    - Most allocations become garbage very soon
    - So focus reclamation work on zones of memory recently allocated
- **For more information:**
  Jones and Lin, "*Garbage Collection: Algorithms for Automatic Dynamic Memory*", John Wiley & Sons, 1996.

Memory as a Graph

- **We view memory as a directed graph**
  - Each block is a node in the graph
  - Each pointer is an edge in the graph
  - Locations not in the heap that contain pointers into the heap are called *root* nodes (e.g. registers, locations on the stack, global variables)

A node (block) is *reachable* if there is a path from any root to that node.

Non-reachable nodes are *garbage* (cannot be needed by the application)

Mark and Sweep Collecting

- **Can build on top of malloc/free package**
  - Allocate using `malloc` until you "run out of space"
- **When out of space:**
  - Use extra *mark bit* in the head of each block
  - *Mark:* Start at roots and set mark bit on each reachable block
  - *Sweep:* Scan all blocks and free blocks that are not marked

root

Before mark

After mark — Mark bit set

After sweep

Note: arrows here denote memory refs, not free list ptrs.

Assumptions For a Simple Implementation

- **Application**
  - **new(n):** returns pointer to new block with all locations cleared
  - **read(b,i):** read location **i** of block **b** into register
  - **write(b,i,v):** write **v** into location **i** of block **b**

- **Each block will have a header word**
  - addressed as **b[-1]**, for a block **b**
  - Used for different purposes in different collectors

- **Instructions used by the Garbage Collector**
  - **is_ptr(p):** determines whether **p** is a pointer
  - **length(b):** returns the length of block **b**, not including the header
  - **get_roots():** returns all the roots

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {
    if (!is_ptr(p)) return;
    if (markBitSet(p)) return;
    setMarkBit(p);
    for (i=0; i < length(p); i++)
      mark(p[i]);
    return;
}
```

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {
   if (!is_ptr(p)) return;        // if not pointer -> do nothing
   if (markBitSet(p)) return;
   setMarkBit(p);
   for (i=0; i < length(p); i++)
     mark(p[i]);
   return;
}
```

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {
   if (!is_ptr(p)) return;        // if not pointer -> do nothing
   if (markBitSet(p)) return;     // if already marked -> do nothing
   setMarkBit(p);
   for (i=0; i < length(p); i++)
     mark(p[i]);
   return;
}
```

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {
   if (!is_ptr(p)) return;        // if not pointer -> do nothing
   if (markBitSet(p)) return;     // if already marked -> do nothing
   setMarkBit(p);                 // set the mark bit
   for (i=0; i < length(p); i++)
     mark(p[i]);
   return;
}
```

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {
   if (!is_ptr(p)) return;        // if not pointer -> do nothing
   if (markBitSet(p)) return;     // if already marked -> do nothing
   setMarkBit(p);                 // set the mark bit
   for (i=0; i < length(p); i++)  // for each word in p's block
     mark(p[i]);
   return;
}
```

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {
   if (!is_ptr(p)) return;          // if not pointer -> do nothing
   if (markBitSet(p)) return;       // if already marked -> do nothing
   setMarkBit(p);                   // set the mark bit
   for (i=0; i < length(p); i++)    // for each word in p's block
     mark(p[i]);                    //  make recursive call
   return;
}
```

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {
   if (!is_ptr(p)) return;       // if not pointer -> do nothing
   if (markBitSet(p)) return;    // if already marked -> do nothing
   setMarkBit(p);                // set the mark bit
   for (i=0; i < length(p); i++) // for each word in p's block
     mark(p[i]);                 //   make recursive call
   return;
}
```

Sweep using lengths to find next block

```
ptr sweep(ptr p, ptr end) {
   while (p < end) {             // for entire heap
      if markBitSet(p)
         clearMarkBit();
      else if (allocateBitSet(p))
         free(p);
      p += length(p);
}
```

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {
   if (!is_ptr(p)) return;        // if not pointer -> do nothing
   if (markBitSet(p)) return;     // if already marked -> do nothing
   setMarkBit(p);                 // set the mark bit
   for (i=0; i < length(p); i++)  // for each word in p's block
     mark(p[i]);                  //  make recursive call
   return;
}
```

Sweep using lengths to find next block

```
ptr sweep(ptr p, ptr end) {
   while (p < end) {              // for entire heap
      if markBitSet(p)           // did we reach this block?
         clearMarkBit();
      else if (allocateBitSet(p))
         free(p);
      p += length(p);
}
```

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {
   if (!is_ptr(p)) return;        // if not pointer -> do nothing
   if (markBitSet(p)) return;     // if already marked -> do nothing
   setMarkBit(p);                 // set the mark bit
   for (i=0; i < length(p); i++)  // for each word in p's block
     mark(p[i]);                  //  make recursive call
   return;
}
```

Sweep using lengths to find next block

```
ptr sweep(ptr p, ptr end) {
   while (p < end) {              // for entire heap
      if markBitSet(p)           // did we reach this block?
         clearMarkBit();         //  yes -> so just clear mark bit
      else if (allocateBitSet(p))
         free(p);
      p += length(p);
}
```

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {
   if (!is_ptr(p)) return;        // if not pointer -> do nothing
   if (markBitSet(p)) return;     // if already marked -> do nothing
   setMarkBit(p);                 // set the mark bit
   for (i=0; i < length(p); i++)  // for each word in p's block
     mark(p[i]);                  //  make recursive call
   return;
}
```

Sweep using lengths to find next block

```
ptr sweep(ptr p, ptr end) {
   while (p < end) {                  // for entire heap
      if markBitSet(p)                // did we reach this block?
         clearMarkBit();              //  yes -> so just clear mark bit
      else if (allocateBitSet(p))     // never reached: is it allocated?
         free(p);
      p += length(p);
}
```

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {
   if (!is_ptr(p)) return;        // if not pointer -> do nothing
   if (markBitSet(p)) return;     // if already marked -> do nothing
   setMarkBit(p);                 // set the mark bit
   for (i=0; i < length(p); i++)  // for each word in p's block
     mark(p[i]);                  //  make recursive call
   return;
}
```

Sweep using lengths to find next block

```
ptr sweep(ptr p, ptr end) {
   while (p < end) {                  // for entire heap
      if markBitSet(p)                // did we reach this block?
         clearMarkBit();              //  yes -> so just clear mark bit
      else if (allocateBitSet(p))     // never reached: is it allocated?
         free(p);                     //  yes -> its garbage, free it
      p += length(p);
}
```

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {
   if (!is_ptr(p)) return;        // if not pointer -> do nothing
   if (markBitSet(p)) return;     // if already marked -> do nothing
   setMarkBit(p);                 // set the mark bit
   for (i=0; i < length(p); i++)  // for each word in p's block
     mark(p[i]);                  //  make recursive call
   return;
}
```

Sweep using lengths to find next block

```
ptr sweep(ptr p, ptr end) {
   while (p < end) {                  // for entire heap
      if markBitSet(p)                // did we reach this block?
         clearMarkBit();              //  yes -> so just clear mark bit
      else if (allocateBitSet(p))     // never reached: is it allocated?
         free(p);                     //  yes -> its garbage, free it
      p += length(p);                 // goto next block
}
```

Conservative Mark & Sweep in C

- ■ **A "conservative garbage collector" for C programs**
  - ▪ **is_ptr()** determines if a word is a pointer by checking if it points to an allocated block of memory
  - ▪ But, in C pointers can point to the middle of a block

- ■ **So how to find the beginning of the block?**
  - ▪ Can use a balanced binary tree to keep track of all allocated blocks (key is start-of-block)
  - ▪ Balanced-tree pointers can be stored in header (use two additional words)

Left: smaller addresses
Right: larger addresses

Today

- **Explicit free lists**
- **Segregated free lists**
- **Garbage collection**
- **Memory-related perils and pitfalls**

Memory-Related Perils and Pitfalls

- **Dereferencing bad pointers**
- **Reading uninitialized memory**
- **Overwriting memory**
- **Referencing nonexistent variables**
- **Freeing blocks multiple times**
- **Referencing freed blocks**
- **Failing to free blocks**

C operators

Postfix

Operators	*Associativity*
`( )  [ ]   ->   .  ++ --`	left to right
`!  ~  ++ --  +  -  *  & (type) sizeof`	right to left
`*  /  %`	left to right
`+  -`	left to right
`<<  >>`	left to right
`<  <=  >  >=`	left to right
`==  !=`	left to right
`&`	left to right
`^`	left to right
`|`	left to right
`&&`	left to right
`||`	left to right
`?:`	right to left
`= += -= *= /= %= &= ^= != <<= >>=`	right to left
`,`	left to right

Unary

Prefix

Binary

- **`->`, `( )`, and `[ ]` have high precedence, with `*` and `&` just below**
- **Unary `+`, `-`, and `*` have higher precedence than binary forms**

Source: K&R page 53, updated

C Pointer Declarations: Test Yourself!

`int *p`   p is a pointer to int

`int *p[13]`   p is an array[13] of pointer to int

`int *(p[13])`   p is an array[13] of pointer to int

`int **p`   p is a pointer to a pointer to an int

`int (*p)[13]`   p is a pointer to an array[13] of int

`int *f()`   f is a function returning a pointer to int

`int (*f)()`   f is a pointer to a function returning int

`int (*(*f())[13])()`   f is a function returning ptr to an array[13] of pointers to functions returning int

`int (*(*x[3])())[5]`   x is an array[3] of pointers to functions returning pointers to array[5] of ints

Source: K&R Sec 5.12

Parsing: `int (*(*f())[13])()`

```
int (*(*f())[13])()        f

int (*(*f())[13])()        f is a function

int (*(*f())[13])()        f is a function
                           that returns a ptr

int (*(*f())[13])()        f is a a function
                           that returns a ptr to an
                           array of 13

int (*(*f())[13])()        f is a ptr to a function
                           that returns a ptr to an
                           array of 13 ptrs

int (*(*f())[13])()        f is a ptr to a function
                           that returns a ptr to an
                           array of 13 ptrs to function
                           returning an int
```

C Pointer Declarations: Test Yourself!

`int *p`   p is a pointer to int

`int *p[13]`   p is an array[13] of pointer to int

`int *(p[13])`   p is an array[13] of pointer to int

`int **p`   p is a pointer to a pointer to an int

`int (*p)[13]`   p is a pointer to an array[13] of int

`int *f()`   f is a function returning a pointer to int

`int (*f)()`   f is a pointer to a function returning int

`int (*(*f())[13])()`   f is a function returning ptr to an array[13]
                                      of pointers to functions returning int

`int (*(*x[3])())[5]`   x is an array[3] of pointers to functions
                                      returning pointers to array[5] of ints

Source: K&R Sec 5.12

A better way: `int (*(*f())[13])()`

// pointer to a function returning an int
```
typedef int (*pfri)();
```

// An array of thirteen pfri's
```
typedef pfri arr13pfri[13];
```

// pointer to an array of thirteen pfri's
```
typedef arr13pfri* ptrToArr;
```

// ptr to function returning a
// ptr to an array of 13 pointer's to functions which return ints
```
typedef ptrToArr (*pfrArr13fri)();
```

Dereferencing Bad Pointers

- ## The classic `scanf` bug

```
int val;

...

scanf("%d", val);
```

Reading Uninitialized Memory

■ **Assuming that heap data is initialized to zero**

```
/* return y = Ax */
int *matvec(int **A, int *x) {
    int *y = malloc(N*sizeof(int));
    int i, j;

    for (i=0; i<N; i++)
        for (j=0; j<N; j++)
            y[i] += A[i][j]*x[j];
    return y;
}
```

■ **Can avoid by using** `calloc`

Overwriting Memory

■ **Allocating the (possibly) wrong sized object**

```
int **p;

p = malloc(N*sizeof(int));

for (i=0; i<N; i++) {
    p[i] = malloc(M*sizeof(int));
}
```

■ **Can you spot the bug?**

Overwriting Memory

- **Off-by-one errors**

```
char **p;

p = malloc(N*sizeof(int *));

for (i=0; i<=N; i++) {
    p[i] = malloc(M*sizeof(int));
}
```

```
char *p;

p = malloc(strlen(s));
strcpy(p,s);
```

Overwriting Memory

- **Not checking the max string size**

```
char s[8];
int i;

gets(s);   /* reads "123456789" from stdin */
```

- **Basis for classic buffer overflow attacks**

Overwriting Memory

- **Misunderstanding pointer arithmetic**

```
int *search(int *p, int val) {

    while (p && *p != val)
        p += sizeof(int);

    return p;
}
```

Overwriting Memory

- **Referencing a pointer instead of the object it points to**

```
int *BinheapDelete(int **binheap, int *size) {
    int *packet;
    packet = binheap[0];
    binheap[0] = binheap[*size - 1];
    *size--;
    Heapify(binheap, *size, 0);
    return(packet);
}
```

Operators												Associativity
()	[]	->	.	++	--							left to right
!	~	++	--	+	-	*	&	(type)	sizeof			right to left
*	/	%										left to right
+	-											left to right
<<	>>											left to right
<	<=	>	>=									left to right
==	!=											left to right
&												left to right
^												left to right
|												left to right
&&												left to right
||												left to right
?:												right to left
=	+=	-=	*=	/=	%=	&=	^=	!=	<<=	>>=		right to left
,												left to right

Referencing Nonexistent Variables

- **Forgetting that local variables disappear when a function returns**

```
int *foo () {
    int val;

    return &val;
}
```

Freeing Blocks Multiple Times

- **Nasty!**

```
x = malloc(N*sizeof(int));
        <manipulate x>
free(x);

y = malloc(M*sizeof(int));
        <manipulate y>
free(x);
```

Referencing Freed Blocks

- **Evil!**

```
x = malloc(N*sizeof(int));
    <manipulate x>
free(x);
    ...
y = malloc(M*sizeof(int));
for (i=0; i<M; i++)
    y[i] = x[i]++;
```

Failing to Free Blocks (Memory Leaks)

- **Slow, long-term killer!**

```
foo() {
    int *x = malloc(N*sizeof(int));
    ...
    return;
}
```

Failing to Free Blocks (Memory Leaks)

- **Freeing only part of a data structure**

```
struct list {
    int val;
    struct list *next;
};

foo() {
    struct list *head = malloc(sizeof(struct list));
    head->val = 0;
    head->next = NULL;
    <create and manipulate the rest of the list>
     ...
    free(head);
    return;
}
```

Dealing With Memory Bugs

- **Debugger: `gdb`**
  - Good for finding bad pointer dereferences
  - Hard to detect the other memory bugs

- **Data structure consistency checker**
  - Runs silently, prints message only on error
  - Use as a probe to zero in on error

- **Binary translator: `valgrind`**
  - Powerful debugging and analysis technique
  - Rewrites text section of executable object file
  - Checks each individual reference at runtime
    - Bad pointers, overwrites, refs outside of allocated block

- **glibc malloc contains checking code**
  - `setenv MALLOC_CHECK_  3`