

# Malloc Boot Camp

Stan, Nikhil, Kim

# Agenda

- Conceptual Overview
  - Explicit List
  - Segregated list
  - Splitting, coalescing
  - Hints on hints
- Advanced debugging with GDB
  - Fun GDB tricks
- Writing a good heap checker
- Appendix

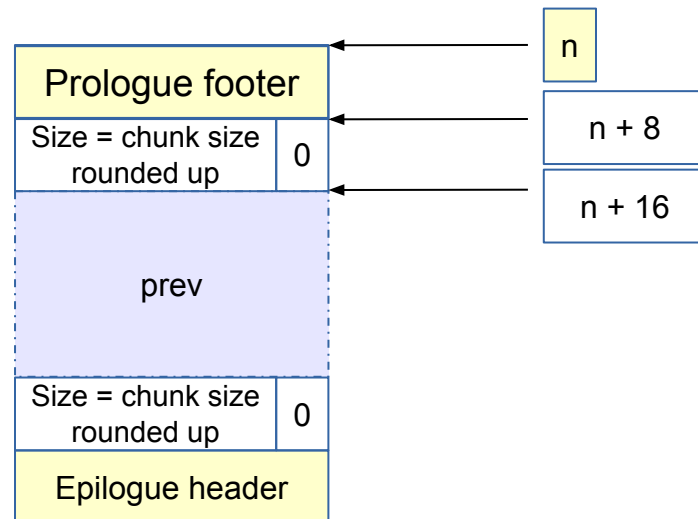

# Conceptual Outline

# Dynamic Memory Allocation

- Used when
  - we don't know at compile-time how much memory we will need
  - when a particular chunk of memory is not needed for the entire run
    - lets us re-use that memory for storing other things
- Important terms:
  - malloc/calloc/realloc/free
  - mem\_sbrk
  - payload
  - fragmentation (external vs internal)
  - Splitting / coalescing

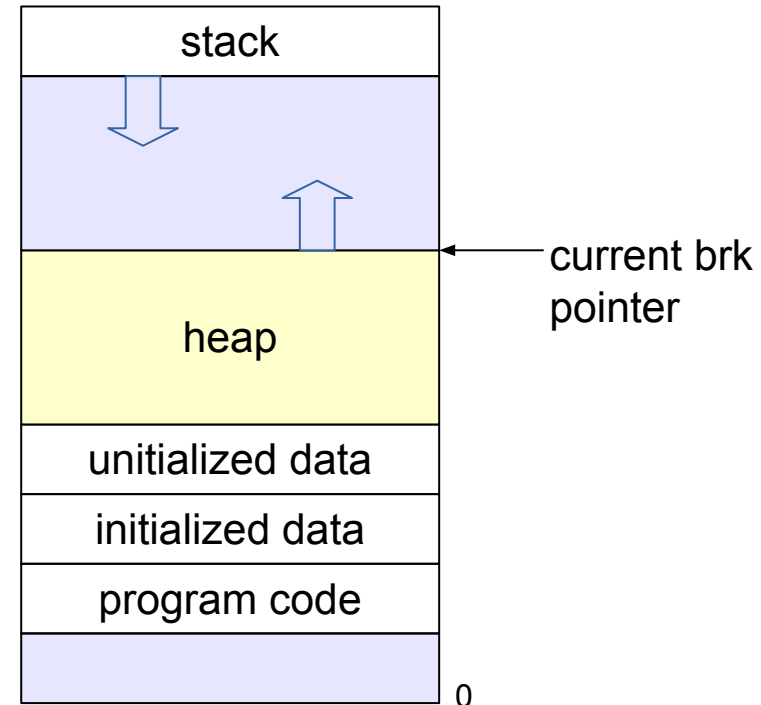
# mm\_init

- Why prologue footer and epilogue header?
- Payload must be 16-byte aligned
- But, the size of payload doesn't have to be a multiple of 16!
- Things malloc'd must be within the prologue and epilogue



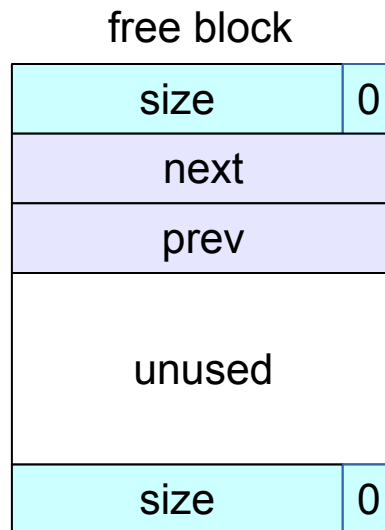

# If We Can't Find a Usable Free Block

- Need to extend the heap
  - in Malloclab, use `mem_sbrk()`
    - `sbrk(requested_bytes)` allocates `requested_bytes` of space and returns pointer to start
    - `sbrk(0)` returns a pointer to the end of the current heap
- For speed, extend the heap by a little more than you need immediately (`chunksize`)
  - use what you need out of the new space, add the rest as a free block




# Tracking Blocks: Explicit List

- Maintain a list of *free* blocks instead of *all* blocks
  - means we need to store forward/backward pointers, not just sizes
  - we only track free blocks, so we can store the pointers in the payload area!
  - need to store size at end of block too, for coalescing



# Splitting a Block


- If the block we find is larger than we need, split it and leave the remainder for a future allocation
  - explicit lists: correct previous and next pointers
  - Segregated lists: same as explicit
- When would we **not** split a block?



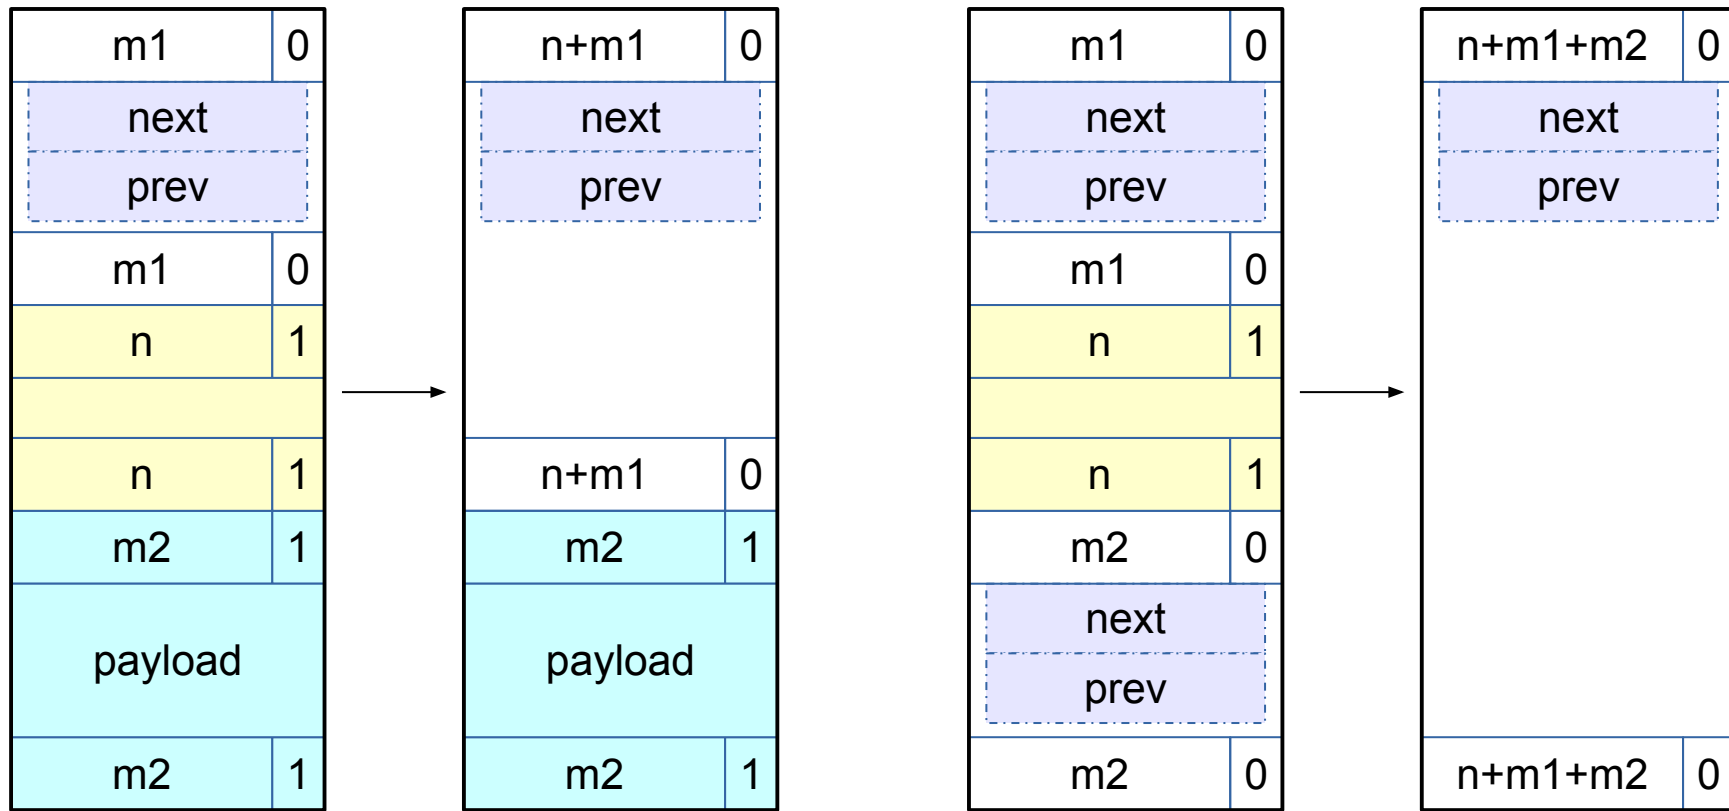


# Coalescing Memory

- Combine adjacent blocks if both are free
  - **explicit lists**: look forward and backward using block sizes, **not** next/prev
  - **footerless**: if free, obtain info from footer then use next/prev
  - **segregated lists**: look forward and back using block sizes, then
    - use the size of the coalesced block to determine the proper list
    - insert into list using the insertion policy (LIFO, address-ordered, etc.)
- Four cases:



# Coalescing Memory



# Design Considerations

- Finding a matching free block
  - First fit vs. next fit vs. best fit
  - continue searching for a closer fit after finding a big-enough free block?
- Free block ordering
  - LIFO, FIFO, or address-ordered?
- When to coalesce
  - while freeing a block or while searching for free memory?
- How much memory to request with `sbrk()`
  - larger requests save time in system calls but increase maximum memory use

# Hints on hints

For the final, you must greatly increase the utilization and keep a high throughput.


- Reducing external fragmentation requires achieving something closer to best-fit allocated e.g., by using segregated lists.
- Reducing internal fragmentation requires reducing data structure overhead and using the a 'good' free block

# Segregated Lists

- Just multiple explicit lists where the free blocks are of a certain size range
- Increases throughput and raises probability of choosing a better-sized block
- Need to decide what size classes and you are bounded by 128 byte allowance
- RootSizeClass1 -> free-block 1 -> free-block 2 -> free-block 3 ->
- RootSizeClass2 -> free-block 1 -> free-block 2 -> free-block 3 -> ...
- RootSizeClass3 -> free-block 1 -> free-block 2 -> free-block 3 -> ...
- ...

# Eliminate footers in allocated blocks

- Reduces internal fragmentation (increase utilization)
- How to coalesce when the allocated block is free?
- Consider how we stored information about whether a block was allocated
- Free blocks still have footers
- Based on that idea, can you design a method that helps you determine when to coalesce?




# Decrease the minimum block size

- Reduces internal fragmentation (increase utilization)
- Currently, minimum block size is 32.
- 8 byte header, 16 byte payload (or 2 8 byte pointers for free), 8 byte footer
- If you just need to malloc(5), and the payload size is 16, you waste 9.
- Must manage free blocks that are too small to hold the pointers for a **doubly** linked free list

m1	1
payload	
m1	1

# Header Reduction

- When to use 8 or 4 byte header? (must support all possible block sizes)
- If 4 byte, how to ensure that payload is aligned?
- Arrange accordingly
- How to coalesce if 4 byte header block is followed by 8 byte header block?
- Store extra information in headers





# Preventing Errors

- Good coding practices can make your code less error-prone
- Plan what each function does before writing it
  - consider edge cases – block at start/end of list, single item on list, etc.
  - draw pictures to help you visualize linked lists, memory layout, etc.
- Document your code as you write it
- Use helper functions!
- Check for common errors:
  - dereferencing invalid pointers / reading uninitialized memory
  - overwriting memory
  - freeing blocks multiple times (or not at all) / referencing freed blocks
  - incorrect pointer arithmetic

# Debugging: GDB & The Almighty Heap Checker

# Download this handout for fun & profit

- Presentation on course website
  - See “schedule”
- <https://cs.cmu.edu/~213/activities/mallocbootcamp.tar>



...except it's not.

# Better than printf: using GDB

- Use GDB to determine where segfaults happen!
- gdb mdriver will open the malloc driver in gdb
  - Type run and your program will run until it hits the segfault!
- layout split - display GDB “text user interface”
  - Way prettier than normal GDB command prompt
  - Can sometimes get messed up by programs that print things
  - refresh - redraws screen
- step - step to the next line of code, stepping into functions if necessary
- next - same as above, but steps *over* functions instead
- finish - continue execution until the end of the current function, then break

# Using GDB - Fun with frames

- backtrace - print call stack up until current function
  - backtrace full - print local variables in each stack frame

(gdb) backtrace

#0 find\_fit (...)

#1 mm\_malloc (...)

#2 0x0000000000403352 in eval\_mm\_valid (...)

#3 run\_tests (...)

#4 0x0000000000403c39 in main (...)

- frame 1 - switch to mm\_malloc's stack frame
  - Good for inspecting local variables of calling functions

# Using GDB - Setting breakpoints/watchpoints

- `break mm_checkheap` - break on function “`mm_checkheap`”
  - `break mm.c:25` - break on line 25 of file “`mm.c`” - **very useful!**
- `break find_fit if size == 24` - break on function “`find_fit`” if the local variable “`size`” in the `malloc` function is equal to 24
  - See `mm-baseline.c` in `malloc` handout
- `watch heap_listp` - break if the value of “`heap_listp`” changes
- `watch block == 0x80000010` - break if “`block`” is equal to this value
- `watch *0x15213` - watch for changes at memory location `0x15213`
  - Can be *very* slow
- `rwatch <thing>` - stop on reading a memory location
- `awatch <thing>` - stop on *any* memory access

# GDB Live Demo



# Heap Checker

- `int mm_checkheap(int verbose);`
- critical for debugging
  - **write this function early!**
  - update it when you change your freelist implementation
  - check all heap invariants (next slide), make sure you haven't lost track of any part of your heap
    - check should pass if and only if the heap is truly well-formed
  - should only generate output if a problem is found, to avoid cluttering up your program's output
- meant to be correct, not efficient
- call before/after major operations when the heap should be well-formed

# Heap Invariants (**Non-Exhaustive**)

- Block level
  - What are some things which should always be true of every block in the heap?

# Heap Invariants (**Non-Exhaustive**)

- Block level
  - header and footer match
  - payload area is aligned, size is valid
  - no contiguous free blocks unless you defer coalescing
- List level
  - What are some things which should always be true of every element of a free list?

# Heap Invariants (**Non-Exhaustive**)

- Block level
  - header and footer match
  - payload area is aligned, size is valid
  - no contiguous free blocks unless you defer coalescing
- List level
  - next/prev pointers in consecutive free blocks are consistent
  - no allocated blocks in free list, all free blocks are in the free list
  - no cycles in free list unless you use a circular list
  - each segregated list contains only blocks in the appropriate size class
- Heap level
  - What are some things that should be true of the heap as a whole?

# Heap Invariants (**Non-Exhaustive**)

- Block level
  - header and footer match
  - payload area is aligned, size is valid
  - no contiguous free blocks unless you defer coalescing
- List level
  - next/prev pointers in consecutive free blocks are consistent
  - no allocated blocks in free list, all free blocks are in the free list
  - no cycles in free list unless you use a circular list
  - each segregated list contains only blocks in the appropriate size class
- Heap level
  - all blocks between heap boundaries, correct sentinel blocks (if used)

# Heap Checker Live Demo

# How to Ask for Help

- Be specific about what the problem is, and how to cause it
  - **BAD**: “My program segfaults.”
  - **GOOD**: “On the third free() in trace 4, I get an invalid pointer in my free list while coalescing memory.”
  - Try to figure out which part of the trace file triggers the problem
  - What sequence of events do you expect around the time of the error? What part of the sequence has already happened?
- Have you written your mm\_checkheap function, and is it working?
  - We **WILL** ask to see it!
- Practice asking your rubber duck about the problem (see Recitation 9) before asking a TA or instructor

# If You Get Stuck

- ***Please read the writeup!***
- CS:APP Chapter 9
- View lecture notes and course FAQ at <http://www.cs.cmu.edu/~213>
- Office hours Sunday through Thursday 5:00-9:00pm in WeH 5207
- Post a **private** question on Piazza
- Obtain a rubber duck....



# APPENDIX

# Internal Fragmentation

- Occurs when the *payload* is smaller than the block size
  - due to alignment requirements
  - due to management overhead
  - as the result of a decision to use a larger-than-necessary block
- Depends on the current allocations, i.e. the pattern of *previous* requests

# Internal Fragmentation

- Due to alignment requirements – the allocator doesn't know how you'll be using the memory, so it has to use the strictest alignment:
  - `void *m1 = malloc(13); void *m2 = malloc(11);`
  - `m1` and `m2` both have to be aligned on 8-byte boundaries
- Due to management overhead (each cell is 2 bytes):

l	e	n	1	p	a	y	l	o	a	d	1	l	e	n	2	p	a	y	l	o	a	d	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---


# External Fragmentation

- Occurs when the total free space is sufficient, but no single free block is large enough to satisfy the request
- Depends on the pattern of *future* requests
  - thus difficult to predict, and any measurement is at best an estimate
- Less critical to malloc traces than internal fragmentation

p5 = malloc(4)

free(p1)

p6 = malloc(5)



**Oops! Seven bytes available, but not in one chunk....**

# C: Pointer Arithmetic

- Adding an integer to a pointer is different from adding two integers
- The value of the integer is always multiplied by the size of the type that the pointer points at
- Example:
  - `type_a *ptr = ...;`
  - `type_a *ptr2 = ptr + a;`
- is really computing
  - `ptr2 = ptr + (a * sizeof(type_a));`
  - i.e. `lea (ptr, a, sizeof(type_a)), ptr2`
- Pointer arithmetic on `void*` is undefined (what's the size of a void?)

# C: Pointer Arithmetic

- `int *ptr = (int*)0x152130;`

- `int *ptr2 = ptr + 1;`

- `char *ptr = (char*)0x152130;`

- `char *ptr2 = ptr + 1;`

- `char *ptr = (char*)0x152130;`

- `void *ptr2 = ptr + 1;`

- `char *ptr = (char*)0x152130;`


- `char *p2 = ((char*)((int*)ptr)+1));`

# C: Pointer Arithmetic


- `int *ptr = (int*)0x152130;`  
`int *ptr2 = ptr + 1; // ptr2 is 0x152134`
- `char *ptr = (char*)0x152130;`  
`char *ptr2 = ptr + 1; // ptr2 is 0x152131`
- `char *ptr = (char*)0x152130;`  
`void *ptr2 = ptr + 1; // ptr2 is still 0x152131`
- `char *ptr = (char*)0x152130;`  
`char *p2 = ((char*)((int*)ptr)+1); // p2 is 0x152134`

# Dynamic Memory Allocation: Example


p1 = malloc(3)




p2 = malloc(7)




p3 = malloc(5)




free(p2)



p4 = malloc(4)



p5 = malloc(4)





# The Memory-Block Information Data Structure


- Requirements:
  - tells us where the blocks are, how big they are, and whether they are free
  - must be able to update the data during calls to malloc and free
  - need to be able to find the **next free block** which is a “good enough fit” for a given payload
  - need to be able to quickly mark a block as free or allocated
  - need to be able to detect when we run out of blocks
    - what do we do in that case?
- The only memory we have is what we're handing out
  - ...but not all of it needs to be payload! We can use part of it to store the block information.

# Finding a Free Block


- First Fit
  - search from beginning, use first block that's big enough
  - linear time in total number of blocks
  - can cause small “splinters” at beginning of list
- Next Fit
  - start search from where previous search finished
  - often faster than first fit, but some research suggests worse fragmentation
- Best Fit
  - search entire list, use smallest block that's big enough
  - keeps fragments small (less wasted memory), but slower than first fit

# Freeing Blocks

- Simplest implementation is just clearing the “allocated” flag
  - but leads to external fragmentation



`free(p)`



`malloc(8)`

*Oops!*

# Insertion Policy


- Where do you put a newly-freed block in the free list?
  - LIFO (last-in-first-out) policy
    - add to the beginning of the free list
    - pro: simple and constant time (very fast)  
*block->next = freelist; freelist = block;*
    - con: studies suggest fragmentation is worse
  - Address-ordered policy
    - insert blocks so that free list blocks are always sorted by address  
 $\text{addr}(\text{prev}) < \text{addr}(\text{curr}) < \text{addr}(\text{next})$
    - pro: lower fragmentation than LIFO
    - con: requires search

# C: Pointer Casting

- Notation:  $(b^*)$  a “casts” a to be of type  $b^*$
- Casting a pointer doesn't change the bits!
  - `type_a *ptr_a=...; type_b *ptr_b=(type_b*)ptr_a;`  
makes `ptr_a` and `ptr_b` contain identical bits
- But it does change the behavior when dereferencing
  - because we *interpret* the bits differently
- Can cast `type_a*` to long/unsigned long and back
  - pointers are really just 64-bit numbers
  - such casts are important for mallocab
  - but be careful – this can easily lead to hard-to-find errors

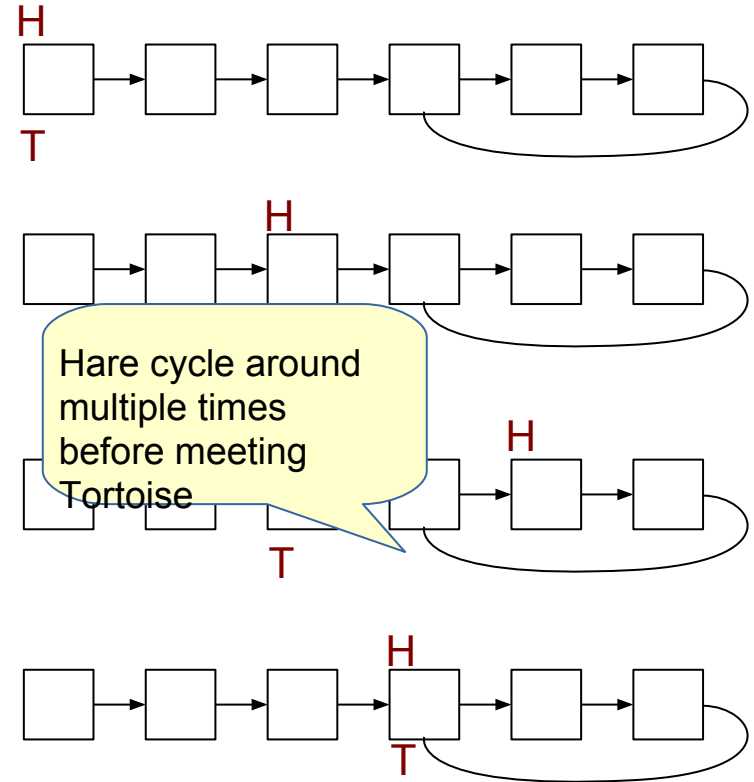
# Cycle Checking: Hare and Tortoise Algorithm

- This algorithm detects cycles in linked lists
- Set two pointers, called “hare” and “tortoise”, to the beginning of the list
- During each iteration, move “hare” forward by two nodes, “tortoise” by one node
  - if “tortoise” reaches the end of the list, there is no cycle
  - if “tortoise” equals “hare”, the list has a cycle



# Cycle Checking: Hare and Tortoise Algorithm

- This algorithm detects cycles in linked lists
- Set two pointers, called “hare” and “tortoise”, to the beginning of the list
- During each iteration, move “hare” forward by two nodes, “tortoise” by one node
  - if “tortoise” reaches the end of the list, there is no cycle
  - if “tortoise” equals “hare”, the list has a cycle



# Debugging Tip: Using the Preprocessor

- Use conditional compilation with `#if` or `#ifdef` to easily turn debugging code on or off

```
#ifdef DEBUG
# define DBG_PRINTF(...) fprintf(stderr, __VA_ARGS__)
# define CHECKHEAP(verbose) mm_checkheap(verbose)
#else
# define DBG_PRINTF(...)
# define CHECKHEAP(verbose)
#endif /* DEBUG */
```

```
void free(void *p)
{
    DBG_PRINTF("freeing %lx\n", (long)p);
    CHECKHEAP(1);
    ...
}
```



# Debugging Tip: GDB

- Use breakpoints / conditional breakpoints
  - `break {address} if {condition}`
- Use watchpoints
  - like breakpoints, but stop the program when the watched expression changes or location is written
  - `watch {expression}`      `watch block->next`
    - break any time the expression changes value; can be extremely slow!
  - `watch -l {expression}`      `watch -l *0x15213`
    - evaluate the expression and watch the memory location at that address
    - program runs at full speed if GDB can set a hardware watchpoint
  - `rwatch` to stop on *reading* a location, `awatch` to stop on **any** access